

Ionian University

Department of Informatics



Undergraduate Thesis

Developing a Mobile-Assisted Language Learning Application Utilizing Short-Form Video Reels and Hypercasual Games

Mohammad Matin Marzie

Registration Number: [REGISTRATION NUMBER]

Supervisor: Prof. [SUPERVISOR NAME]

September 16, 2026

Ιόνιο Πανεπιστήμιο

Τμήμα Πληροφορικής



Πτυχιακή Εργασία

**Ανάπτυξη Εφαρμογής για Φορητές Συσκευές για τη
Γλωσσική Εκμάθηση αξιοποιώντας Βίντεο Μικρού
Μήκους και Hypercasual Παιχνιδιών**

Mohammad Matin Marzie

Αριθμός Μητρώου: [REGISTRATION NUMBER]

Επιβλέπων: Καθ. [SUPERVISOR NAME]

16 Σεπτεμβρίου 2026

Supervisor

[SUPERVISOR NAME], *Professor,*
Ionian University

Examination Committee

[SUPERVISOR NAME], *Professor,*
Ionian University

[COMMITTEE MEMBER 2 NAME], *Associate Professor,*
Ionian University

[COMMITTEE MEMBER 3 NAME], *Assistant Professor,*
Ionian University

To my family.

All copyrights are reserved by the author.

Abstract

Language Learning has become ...

Key words: [KEYWORD-1], [KEYWORD-2], [KEYWORD-3]

This thesis proposes a Personalized Language Learning (PLL) application that...

Περίληψη

Η παρούσα διπλωματική εργασία ...

Λέξεις κλειδιά: [ΛΕΞΕΙΣ-ΚΛΕΙΔΙΑ], [KEYWORD-1], [KEYWORD-2], [KEYWORD-3]

Acknowledgments

I would like to warmly thank everyone who contributed to the completion of this thesis.

First and foremost, I thank my supervisor Professor [SUPERVISOR NAME] for the guidance, patience and valuable help provided throughout the development of this work. His/her advice and suggestions were crucial for the completion of this project.

I also thank the members of the examination committee for the time they devoted to evaluating my work and for their constructive comments.

I would like to thank all the professors of the Department of Informatics who contributed to my education during my undergraduate studies.

Special thanks go to my family for their unwavering support, understanding and encouragement that they provided me throughout all these years of studies.

Finally, I thank my classmates and friends for their support and the beautiful moments we shared together.

Corfu, [MONTH] 2026

Mohammad Matin Marzie

Contents

A	Introduction	1
A.1	Motivation	1
A.2	Background	2
A.2.1	Short-form Video	2
A.2.2	Hypercasual and Hybrid-Casual Games	3
A.3	Problem Statement and Research Questions	4
A.4	Significance	5
A.5	Thesis Structure	5
B	Literature Review	6
B.1	The Landscape of Mobile-Assisted Language Learning (MALL)	6
B.1.1	Market Evolution and Platform Proliferation	6
B.1.2	Critical Review and Systematic Limitations of Current Platforms	7
B.2	Pedagogical Foundations	9
B.2.1	Sociocultural Theory and the Zone of Proximal Development (ZPD)	10
B.2.2	The Input Hypothesis and the Lexical Coverage Threshold	10
B.2.3	Cognitive Load Theory and Microlearning	11
B.2.4	Memory Decay and Spaced Repetition Models	11
B.2.4.1	<i>The Ebbinghaus Forgetting Curve</i>	11
B.2.4.2	<i>Classic Leitner's Algorithm</i>	12
B.2.4.3	<i>Parametric Half-Life Regression</i>	13
B.2.4.4	<i>FSRS and the DHP Model</i>	13
B.2.5	Active Recall and the Testing Effect	15
B.3	Modalities for Language Acquisition	16

B.3.1	Short-Form Video in Second Language Acquisition (SLA)	16
B.3.1.1	<i>Incidental Vocabulary Acquisition through Audiovisual Input</i>	16
B.3.1.2	<i>The Gap: Passive Consumption without Reinforcement</i> . . .	17
B.3.2	Serious Games and Gamification in Language Learning	17
B.3.2.1	<i>Game-Based Learning versus Gamification</i>	17
B.3.2.2	<i>Difficulty Adaptation in Hypercasual Games</i>	18
B.3.2.3	<i>Active-Recall Game Mechanics for Vocabulary</i>	18
B.4	Algorithmic Frameworks for Personalized Recommendation	18
B.4.1	Content-Based Filtering	19
B.4.2	Collaborative Filtering and Hybrid Recommender Systems	19
B.4.3	Recommender Systems in Educational Technology	19
B.5	Synthesis: A Theory-Driven Design Framework	20

C System Architecture & Methodology 22

C.1	System Overview	22
C.2	Technology Stack	23
C.2.1	Frontend — React Native / Expo	23
C.2.2	Backend API — Node.js / Express	24
C.2.3	Reels Service — Python / FastAPI	24
C.2.4	Database — PostgreSQL	25
C.3	Data Modelling	25
C.3.1	Linguistic Content	25
C.3.2	User State	27
C.3.3	Engagement Signals	27
C.4	Recommendation Engine Design	27
C.4.1	Stage 1: Comprehensibility Filtering	27
C.4.2	Stage 2: Spaced-Repetition Prioritisation	28
C.4.3	Stage 3: Content-Based Personalisation	28
C.4.4	Combining the Three Stages: A Cascade, Not a Weighted Sum	29
C.5	Gamification Design	29
C.5.1	Progression Currencies: Coins, Energy, and Experience	29
C.5.2	Design Principle: Every Game Round Doubles as an FSRS Review . .	30
C.6	Vocabulary-Driven Games	31

C.6.1	Wordle-Style Game	31
C.6.2	Word of Wonders (Crossword)	31

D Implementation Details 33

D.1	Database Schema	33
D.1.1	Linguistic Content Tables	33
D.1.2	User State Tables	34
D.1.3	Engagement Signal Tables	34
D.2	Backend API	34
D.2.1	Authentication and Authorization	35
	<i>D.2.1.1 JWT Access/Refresh Token Rotation</i>	35
	<i>D.2.1.2 Google OAuth 2.0</i>	35
D.2.2	Registration and Vocabulary Initialisation	36
	<i>D.2.2.1 Proficiency-Based Vocabulary Auto-Seeding</i>	36
D.2.3	Offline Synchronisation	36
	<i>D.2.3.1 Change-Set Ordering: Deletes, Updates, Inserts</i>	37
D.2.4	Dictionary and Letters Modules	37
D.3	Frontend Architecture	37
D.3.1	Navigation and Screen Structure	37
D.3.2	State Management: Context and Hooks	38
D.3.3	Offline-First Strategy	38
D.4	Reels Service	39
D.4.1	Service Architecture	39
D.4.2	Reel Creation and Dialogue Enrichment	39
D.4.3	Recommendation Engine Implementation	40
	<i>D.4.3.1 Stage 1: Candidate Loading and Comprehensibility Scoring</i>	40
	<i>D.4.3.2 Stage 2: Frontend-Supplied Due-Word Partitioning</i>	40
	<i>D.4.3.3 Stage 3: TF-IDF Profile Vectors and Cosine Similarity</i>	40
D.4.4	Personalisation Implementation Status	41
D.5	Game Modules	42
D.5.1	Wordle-Style Game Implementation	42
D.5.2	Word of Wonders: Procedural Crossword Generation	42
	<i>D.5.2.1 Word Placement and Intersection Search</i>	42

D.5.2.2	<i>Letter-Budget and Word-Priority Heuristics</i>	44
D.5.2.3	<i>Due-Word Injection as the Practice Word</i>	44
E	Evaluation & Results	46
F	Conclusions & Future work	47
F.1	Contributions	47
F.2	Limitations	47
F.3	Future Work	47
	Bibliography	49
	Abbreviations	54
	Glossary	55

List of Figures

B.1	The forgetting curve, with original data from Ebbinghaus [1], as replotted by Murre and Dros [2]. Image: Jaap M. J. Murre, Wikimedia Commons, CC BY 4.0. . . .	11
B.2	The classic five-box Leitner scheduling system. Image: Zirguezi, Wikimedia Commons [3], CC0 1.0.	12
C.3	High-level system architecture: the mobile frontend communicates with the Backend API (<code>api.glosy.gr</code>) and the Reels Service (<code>recs.glosy.gr</code>) over authenticated REST endpoints; both services share a single PostgreSQL instance. .	23
C.4	Entity-relationship diagram of the 21-table schema, grouped into the three axes of Section C.3: Linguistic Content, User State, and Engagement Signals.	26

List of Tables

<i>B.1 Comparative feature matrix: Established LLPs vs. Glosy framework.</i>	9
<i>B.2 Theory-driven design traceability matrix: theoretical foundation, design implication, and the corresponding Glosy component.</i>	20

Chapter A

Introduction

A.1 Motivation

The scale of foreign language acquisition is an unprecedented phenomenon in contemporary applied linguistics. In our globalised society, over 1.5 billion individuals are actively studying a foreign language [4], seeking to facilitate cross-cultural communication, academic mobility, and access to global scientific literature. This widespread interest in language learning in conjunction with the advent of technological advancement (e.g., smartphones, AI) has shifted the paradigm of second language acquisition (SLA) away from traditional classroom-bound instruction toward self-directed, digital learning environments. Currently, independent, mobile-assisted platforms represent the most widespread medium for independent language study, enabling millions of learners to engage with language materials on demand [5].

Despite the ubiquity of mobile-first access and digital resources, contemporary Mobile-Assisted Language Learning (MALL) faces critical pedagogical efficacy bottlenecks. While self-learning mobile applications have become the standard interface for independent L2 (second language) acquisition [5], their long-term educational utility is severely compromised by learner attrition. Industrial Analysis of self-directed language learning applications reveals a severe pedagogical retention deficit, with first-week user attrition rates exceeding 60% across many major platforms [5].

Furthermore, modern lifestyle demands and cognitive workloads compress the time available for structured, long-form study [6]. Traditional, exhaustive learning structures fail to integrate into the daily routines of adult learners, as evidenced by empirical findings indicating that 53% of active

adult learners and half of educational directors identify excessive daily workloads as the primary barrier to ongoing instruction [6]. To circumvent these temporal constraints, learner behaviors have increasingly shifted toward autonomous, bit-sized micro-learning formats.

While modern mobile applications have successfully popularised the micro-learning format, current digital implementations remain heavily constrained by structural and pedagogical bottlenecks. As analyzed in depth within the literature review (Section B), existing MALL platforms frequently rely on uniform, static curricula that fail to adapt to individual learner demographics, lack authentic multimedia integration, and reduce gamification to repetitive, decontextualized translation drills—a phenomenon termed “shallow gamification.” Consequently, these platforms fail to support advanced proficiency development or sustain long-term cognitive engagement. This systemic deficit shifts the scientific inquiry: *how can mobile language learning frameworks transcend these repetitive, drill-based limitations to deliver adaptive, culturally authentic, and cognitively immersive media that aligns with the organic attention habits of modern learners?* Addressing this question requires investigating informal digital environments where contemporary learners willingly and habitually direct their attention.

A.2 Background

To address these completion and engagement deficits, modern Mobile-Assisted Language Learning (MALL) must align with the media consumption realities of contemporary learners. This thesis explores the integration of two prominent digital phenomena: short-form video and hypercasual game mechanics.

A.2.1 Short-form Video

Short-form video is a media format native to contemporary social media platforms, characterized by its brevity—typically ranging from 15 seconds to 3 minutes—and optimized for rapid, algorithmically driven consumption.

Short-form video has become the dominant interface of digital attention across social networks. Statistically, YouTube Shorts has achieved over 200 billion daily views [7], While Meta Platform announced: Video is a particular bright spot, with video time spent on Instagram up more than 30% since last year and 5% more time spent on Facebook [8]. TikTok’s global active advertising audience has reached 1.59 billion adults, representing 27.5% of the global population aged 18 and

older [9].

This media format captures a disproportionate share of mobile attention. Globally, TikTok Android users average 1 hour and 37 minutes per day in the application, Youtube 1H and 25M, Instagram 1H and 13M and Facebook 1H and 7M [10].

This behavioral pattern is heavily concentrated among younger cohorts: 75% of US teens visit YouTube daily, and 61% engage with TikTok daily [11]. Among US adults, daily TikTok usage exhibits a sharp ten-to-one generational gap, captured by 50% daily usage among 18-to-29-year-olds compared to just 5% among those aged 65 and older [12].

Pedagogically, the cognitive appeal of short-form video lies in its brevity and high visual density. While platforms keep raising video length caps (with TikTok allowing up to 10 minutes and Reels up to 3 minutes) [13], consumer behavior remains heavily anchored around rapid completion. Wistia's benchmark analysis of 13 million videos reveals that videos under 60 seconds hold a 52% average engagement rate, meaning viewers watch more than half the runtime before scrolling [14]. In contrast, longer videos see completion rates decay exponentially [14, 15]. Short-form videos therefore represent a highly engaging, authentic delivery mechanism for Second Language Acquisition (SLA).

A.2.2 Hypercasual and Hybrid-Casual Games

Hypercasual video games are defined as highly accessible, mobile-native software products characterized by minimalist user interfaces, instantaneous "tap-to-play" mechanics, and extremely short core gameplay loops [16]. Unlike casual games, which require some level of long-term planning, tutorial acquisition, or system familiarity, hyper-casual games demand no pre-existing cognitive schema or learning curve from the user, relying on immediately intuitive controls (core game mechanics) such as direction and drawing, matching, puzzle, resign and falling, swerving, swiping, tapping/timing and word and board [16, 17]. In recent years, this genre has evolved into "hybrid-casual" gaming, which retains the low-barrier entry and core mechanics of hypercasual games while strategically incorporating mid-core elements, such as progressive meta-features, item collection, and light narrative layer systems [17].

The exceptional appeal and potential addictiveness of these games are central to their design and rapid market expansion [17]. Reflecting this widespread appeal, the hyper-casual and hybrid-casual genres have experienced exponential audience expansion globally, emerging as a dominant leisure activity on both major mobile platforms [18]. The addictiveness of these games are histor-

ically highlighted. Flappy Bird, which was voluntarily withdrawn from download by its creator due to concerns over its addictive and habit-forming impact on players [17].

Furthermore, because they are playable primarily on ubiquitous mobile devices and utilize intuitive user interfaces, hypercasual games reach a broad demographic spectrum. Notably, they appeal strongly to populations that historically do not identify as gamers or have been marginalized by the "male gamer stereotype" [19]. These include female players—who comprise 63% of the mobile gaming market [20]—as well as older adults (aged 65+) seeking cognitive maintenance, mental health preservation, and familial bonding [21].

Despite this high volume of daily engagement, players primarily consume these games as simple digital leisure pursuits to pass time, alleviate boredom, or fill transitional moments during daily routines and commutes [16, 17]. The highly engaging, low-friction mechanics of hypercasual and hybrid-casual games represent a powerful, untapped attention-delivery system. While their mass appeal stems precisely also from this cognitively empty activities, ... TODO: Add here a quesiton, read yang2020hyper maybe to get help.

A.3 Problem Statement and Research Questions

People play a lot Hypercasual games, can they be leveraged to learn a language? How can we turn reels and games into language learning and make it effective? We have to track each word.(personalize) We have to Engineer a Reel recommendation system(algorithm) with personalized words. Personalize each Hypercasual game level.

The central research problem is the absence of a language learning system that simultaneously leverages (i) the authentic contextual richness of user-generated short-form video and (ii) gamified reinforcement adapted to the individual learner's vocabulary state.

Research questions:

- RQ1.** How can a mobile application integrate authentic short-form video content with a word-level vocabulary tracking model to support second language acquisition?
- RQ2.** How should a content recommendation engine be designed to select reels that are both comprehensible and lexically challenging for a given learner, based on forgetting curves?
- RQ3.** To what extent can hypercasual game mechanics, when parameterised by the learner's vo-

cabulary state, serve as effective spaced-repetition exercises within a language learning application?

The question is: Can't i learn language by watching Reels from existing social media platforms? The answer is in next section B.3.1.

This thesis proposes **Glosy**, a cross-platform mobile application that addresses the identified gaps by embedding vocabulary acquisition within short-form authentic video content and reinforcing it through personalised hypercasual games. The system tracks each learner's vocabulary state at the word level, uses that state to select contextually appropriate reels, and adapts game difficulty to the same vocabulary model.

A.4 Significance

Helps learners use their daily habits for productive learning. The application is mobile-first. Supports Offline game play and practice. 3/4 use offline mode [22].

A.5 Thesis Structure

The remainder of this thesis is organised as follows. Chapter B reviews mobile-assisted language learning platforms, identifies their limitations, and establishes the pedagogical framework. Chapter C presents the system architecture, technology stack, data model, and recommendation engine design. Chapter ?? documents the implementation of each component with reference to the source code and database schema. Chapter E evaluates the system against its design objectives and discusses performance. Chapter ?? states the contributions, acknowledges limitations, and proposes future work.

Chapter B

Literature Review

B.1 The Landscape of Mobile-Assisted Language Learning (MALL)

B.1.1 Market Evolution and Platform Proliferation

Digital language learning platforms (LLPs) have become dominant in foreign language education, with over 70 major commercial applications serving millions of active learners through Mobile-Assisted Language Learning (MALL) [23]. The scale attained by the market leaders illustrates the extent of this concentration: Duolingo alone reports 58.7 million daily active users and 140.6 million monthly active users as of mid-2026 [24]. A single platform’s design choices therefore shape the learning experience of a substantial share of the world’s self-directed language learners, which makes a critical examination of the dominant design paradigm — rather than only of usage statistics — a necessary starting point for this thesis.

This rapid expansion has been further accelerated by Generative Artificial Intelligence (AI) and Large Language Models (LLMs), which have streamlined content authoring and curriculum scaling. While authoring the first 100 language courses on Duolingo required over 12 years of manual pedagogical design, generative models facilitated the deployment of 148 additional courses in early 2025 alone [25]. Content authoring is consequently no longer the primary bottleneck to platform growth; as Section B.1.2 argues, the bottleneck has shifted to *pedagogical design* — how the abundant, cheaply-generated content is sequenced, contextualised, and reinforced for the individual learner.

B.1.2 Critical Review and Systematic Limitations of Current Platforms

Karasimos [23] conducted a cross-platform review of five major commercial LLPs — Duolingo, Memrise, Busuu, LingQ, and Babbel — evaluated against a fixed set of pedagogical and usability criteria. Although several of these platforms have iterated substantially since that 2022 review, the underlying architectural pattern it identifies — a static exercise bank gamified with points, streaks, and leaderboards — remains the dominant paradigm in 2026. Building on that review and on our own examination of the current (2026) state of these five platforms, six systematic limitations recur across the category:

1. **Uniform Content Delivery Across Learner Demographics and Interests:** Existing platforms deliver static, uniform curricula that do not account for variations in age cohorts (e.g., Gen Z, adults, seniors), communicative objectives, or cultural nuances. Nor do they account for variation in individual *interest*: two learners with identical proficiency levels but distinct interests — one following football, the other following technology news — are shown the same fixed exercise bank rather than content drawn from the topics that would sustain their engagement. Machine learning implementations in mainstream platforms focus predominantly on tuning exercise *difficulty* rather than personalising semantic *content* to either dimension [22]; that is, adaptivity is applied to *how hard* an item is, not to *who it is for* or *what it is about*.
2. **Insufficient Multimedia Integration:** Primary instructional routines rely predominantly on isolated text prompts and synthetic text-to-speech audio drills. This structure underutilises the visual and paralinguistic memory channels — gesture, facial expression, intonation, situational context — that reinforce natural conversational patterns and prosody, and that a short video clip conveys for free alongside the spoken utterance.
3. **Absence of Authentic, Real-World Linguistic Context:** Authentic input — language produced by native speakers for native speakers within a genuine communicative context, rather than written or graded for pedagogical consumption — is widely regarded as essential for developing communicative competence, and its use is associated with increased learner motivation and cultural awareness [26]. Mainstream LLP curricula, by contrast, frequently isolate target vocabulary into synthetic, decontextualised sentences authored specifically for the exercise. Authentic materials are not a costless substitute, however: they routinely presuppose cultural background knowledge that a foreign learner does not yet possess, which can

impede comprehension [27]. Existing LLPs offer no mechanism through which a confused learner can resolve this kind of gap in situ (Section B.1.2, item 6, returns to this point).

4. **Shallow Gamification:** Nearly all mainstream platforms are *gamified*, in the sense of layering points, streaks, badges, and leaderboards — game-design elements — onto an otherwise conventional drill [28]. Few, however, qualify as *game-based learning*, in which the learning content is embedded within the core mechanical loop of an actual game rather than merely rewarded by one [29]. Duolingo’s exercises, for instance, remain multiple-choice or sentence-construction drills regardless of how many streak-freeze tokens or leaderboard leagues surround them: the reward layer is decorative rather than structural. This distinction matters pedagogically, because points and streaks primarily sustain *extrinsic* engagement (returning to the app), whereas a well-designed game mechanic can additionally scaffold the *cognitive act of retrieval itself*. Section B.3.2 develops this distinction as a pedagogical theory; Chapter C operationalises it in the design of Glosy’s vocabulary-driven games.
5. **Limited Advanced Proficiency Support (B2+):** Current mobile platforms focus heavily on A1–B1 grammar foundations, offering sparse resources for upper-intermediate and advanced learners requiring idiomatic fluency, register variation, and natural speech rates. Except for LingQ, which provides authentic content, most platforms fail to support learners meaningfully beyond the B1 level [23].
6. **Absence of Sociocultural Grounding:** Vygotsky’s sociocultural account of cognitive development holds that learning is fundamentally a socially mediated process, not a solitary one (Section B.2.1) [30, 31]. Current mainstream platforms, however, are architected predominantly around single-player, drill-and-feedback loops with limited collaborative scaffolding: a learner who does not understand an idiom or a cultural reference has no in-context avenue to ask a native speaker, a peer, or an instructor. Busuu is a partial exception among the five platforms reviewed here, offering a community feature in which native speakers correct submitted writing and speech samples; this is a genuine sociocultural mechanism, but it is decoupled from the platform’s own exercises rather than attached to the specific piece of content the learner is struggling with. Outside the five platforms formally reviewed in this section, dedicated language-exchange applications such as HelloTalk [32] address the social-scaffolding gap directly, connecting learners in real time with native speakers for authentic conversation and correction. This is precisely the gap that authentic materials expose in item 3 above — authentic content increases the *need* for social scaffolding at exactly the

point where a structured, vocabulary-tracking curriculum platform provides the *least* of it — and it is the gap *Glosy* targets by pairing a tracked vocabulary model with comment-driven, content-attached social scaffolding, identified as future work in Chapter F.

Table B.1 summarises the resulting comparative feature matrix, contrasting the five established platforms reviewed above with the design targeted by *Glosy*.

Table B.1: *Comparative feature matrix: Established LLPs vs. Glosy framework.*

Feature / Dimension	Duolingo	Memrise	Busuu	Babbel	LingQ	Glosy
Personalized / Interest-Driven Content	Low	Low	Medium	Medium	High	High
Authentic Multimedia Video	Low	Medium	Medium	Low	High	High
Deep Gameplay Mechanics	Low	Low	Low	Low	None	High
B2+ Proficiency Support	Low	Low	Medium	Low	High	High
Adaptive SRS Tracking	High	High	Medium	Medium	Medium	High
Sociocultural / Social Scaffolding	Low	Low	Medium	Low	Medium	High
Open-source	No	No	No	No	No	Yes
Novice (A0–A1) Onboarding Ease	High	High	High	High	Low	Low

The remainder of this chapter builds the theoretical foundation from which the *Glosy* design responds to each of these six limitations: Section B.2 establishes the pedagogical theory, Section B.3 reviews the two content modalities (short-form video and serious games) through which that theory is delivered, and Section B.4 reviews the algorithmic frameworks that select and sequence content for the individual learner. Section B.5 closes the chapter by mapping each theoretical foundation explicitly onto a corresponding design decision realised in Chapter C.

B.2 Pedagogical Foundations

This thesis follows a *theory-driven design* approach: every major design decision documented in Chapter C is traceable to one of the theoretical foundations established in this section, rather than to intuition alone. This section is deliberately scoped to *theory* — what the literature says about how second language vocabulary is acquired, retained, and forgotten — and defers the specific computational mechanisms by which *Glosy* operationalises each theory to Chapter C, which references back to the relevant subsection below at each design decision.

B.2.1 Sociocultural Theory and the Zone of Proximal Development (ZPD)

Vygotsky's Sociocultural Theory posits that cognitive development occurs through socially mediated communication and cultural artefacts, wherein language functions as the primary cognitive tool for organising thought [30]. Central to this paradigm is the **Zone of Proximal Development (ZPD)**: the gap between a learner's independent competence — what they can accomplish unaided — and their potential development when assisted by a more capable interlocutor or by structured scaffolding. Lantolf's application of sociocultural theory to second language acquisition (SLA) re-frames the ZPD specifically for language learning: instructional stimuli should be pitched neither so far below the learner's current competence that no new learning occurs, nor so far above it that the material is incomprehensible without assistance [31]. Instructional design is optimal, in this account, precisely at the boundary of what the learner can process with appropriate support.

The ZPD is a qualitative, dyadic construct: it presupposes a more knowledgeable other (a teacher, a peer, a native speaker) providing scaffolding in real time. Section B.2.2 introduces a complementary, quantitative account of the same boundary — Krashen's Input Hypothesis — which *Glosy* operationalises computationally in the absence of a live interlocutor (Chapter C, Section C.4). The sociocultural dimension of the ZPD — social scaffolding through peer or native-speaker interaction — is not currently implemented in *Glosy*, and is identified as future work in Chapter F in direct response to the gap identified in Section B.1.2, item 6.

B.2.2 The Input Hypothesis and the Lexical Coverage Threshold

Krashen's Input Hypothesis proposes that language acquisition occurs when a learner is exposed to *comprehensible input* slightly beyond their current competence, denoted $i + 1$, where i represents the learner's current linguistic level [33]. Input that consists entirely of already-known language (i) offers no acquisitional value, while input too far beyond the learner's level ($i + n$, $n \gg 1$) is unparseable and therefore equally unproductive. Although formulated independently of Vygotsky, the Input Hypothesis is widely read alongside sociocultural theory in the SLA literature as a compatible, learner-internal account of the same comprehensibility boundary that the ZPD describes at the level of social interaction [31].

Unlike the ZPD, the $i + 1$ boundary admits an empirical, corpus-based operationalisation at the lexical level. Hu and Nation's corpus-based study of unsimplified text established that a reader requires knowledge of approximately 98% of the running word tokens in a text to comprehend it

unaided, with comprehension degrading sharply below that threshold [34]. This figure was subsequently popularised and synthesised as a general vocabulary-learning target in Nation’s widely-cited *Learning Vocabulary in Another Language* [35], and it remains one of the few points at which the largely qualitative construct of comprehensible input admits a precise, machine-checkable definition. This is the specific theoretical result that Chapter C, Section C.4 (Stage 1) operationalises as the comprehensibility filter applied to candidate reels.

B.2.3 Cognitive Load Theory and Microlearning

Cognitive Load Theory (CLT) holds that working memory capacity is strictly bounded — classically estimated at 7 ± 2 discrete chunks of information [36] — and that instructional design must minimise *extraneous* cognitive load (load imposed by poor presentation, unrelated to the material itself) in order to leave sufficient working-memory capacity for *germane* load, the productive cognitive work of constructing new long-term memory schemata [37]. Applied to instructional sequencing, CLT implies that a single learning episode should be bounded in scope: introducing too many new lexical items, grammatical structures, or task demands within one episode exceeds working-memory capacity and displaces schema construction with the effort of merely holding the material in mind.

Microlearning — the pedagogical strategy of decomposing instruction into short, self-contained units rather than long-form sessions — is a direct instructional response to this bound: by construction, a short unit cannot introduce more new material than the working-memory ceiling permits. This principle underlies the length constraints Chapter C imposes on both the video content and the game sessions in *Glosy*: each unit of engagement is scoped narrowly enough that the germane load it imposes remains tractable, consistent with the CLT account above.

B.2.4 Memory Decay and Spaced Repetition Models

B.2.4.1 The Ebbinghaus Forgetting Curve

Ebbinghaus’s foundational experiments on nonsense-syllable recall demonstrated that memory retention decays over time in the absence of reinforcement, most rapidly immediately after initial exposure and more slowly




Figure B.1: The forgetting curve, with original data from Ebbinghaus [1], as replotted by Murre and Dros [2]. Image: Jaap M. J. Murre, Wikimedia Commons, CC BY 4.0.

thereafter [1]. Although Ebbinghaus’s own savings-method data did not yield a single closed-form expression, later reconstructions commonly approximate the resulting *forgetting curve* as a negative exponential function of elapsed time t and a memory-strength parameter S : retention $R(t) = e^{-t/S}$, where a larger S corresponds to a more durably consolidated memory. This exponential-decay account is the shared theoretical ancestor of every spaced-repetition model reviewed below: each attempts, with increasing sophistication, to estimate the parameters of an individual learner’s own decay curve for an individual item, and to schedule review immediately before predicted recall failure.

B.2.4.2 Classic Leitner’s Algorithm

The Leitner system operationalised the forgetting curve as a simple, non-parametric scheduling heuristic decades before it became computationally tractable to fit an individual decay curve per item [38]. Flashcards are sorted into a series of boxes; a card answered correctly is promoted one box forward, to be reviewed at a longer interval, while a card answered incorrectly is reset all the way back to the first box, regardless of how far it had previously advanced, to be reviewed again at the shortest interval. The scheme requires no numerical memory model at all — the box number is a coarse, discretised proxy for the item’s memory strength S — which makes it trivially implementable but insensitive to the fact that different items, and different learners, decay at markedly different rates even within the same box.




Figure B.2: The classic five-box Leitner scheduling system. Image: Zirquezi, Wikimedia Commons [3], CC0 1.0.

B.2.4.3 Parametric Half-Life Regression

Duolingo’s half-life regression (HLR) model, introduced in 2016 [39], replaces the Leitner system’s discrete boxes with a continuous, per-item, per-learner estimate of the memory half-life h — the elapsed time at which recall probability has decayed to 50%. Recall probability p for a given item is modelled as an exponential decay in the elapsed interval Δt since it was last reviewed, directly generalising the Ebbinghaus curve of Section B.2.4.1:

$$p = 2^{-\Delta t/h} \quad (\text{B.1})$$

The half-life h itself is not fixed but estimated per item and per learner via a log-linear regression over a feature vector $\mathbf{x}$ — such as the number of prior correct and incorrect attempts — with a learned weight vector Θ :

$$\hat{h}_{\Theta}(\mathbf{x}) = 2^{\Theta \cdot \mathbf{x}} \quad (\text{B.2})$$

generalising the Leitner box index into a continuously-valued, learner- and item-specific quantity. The weight vector Θ is not known in advance: it must itself be fitted offline, by regression over a large historical corpus of review outcomes pooled across many users, before Equation B.2 can be applied to a single new one — Settles and Meeder trained their original model on Duolingo’s own multi-million-user interaction log [39]. A newly-launched application, without an equivalent corpus of prior review data from which to fit Θ , cannot reproduce this regression from its own history alone. Section B.2.4.4 reviews an alternative memory model that avoids this cold-start requirement.

B.2.4.4 FSRs and the DHP Model

The Free Spaced Repetition Scheduler (FSRS) originates in the DHP model developed at MaiMemo [40], a variant of what the open-source scheduling community terms the **DSR** model: memory state is represented not by a single half-life h , as in HLR, but by three jointly-updated variables — **difficulty** $D \in [1, 10]$ (how intrinsically hard the item is to learn), **stability** S (the number of days for retrievability to decay to 90%), and **retrievability** R (the instantaneous, momentary probability of successful recall, analogous to p in HLR). Each is revised after every individual review rather than re-estimated from scratch, and each review is scored on a four-point grade $G \in \{1, 2, 3, 4\}$ (*again, hard, good, easy*). Since its first release, FSRS has been refined through several major versions (v1 through the current FSRS-6), each adding parameters and improving the empirical fit of its

forgetting curve and update rules [41]; the formulation below is the current one (FSRS-6), noting where it revises earlier versions.

The forgetting curve. Rather than the fixed exponential decay of Equation B.1, FSRS models retrievability as a power-law function of the elapsed interval t and the current stability S , with a trainable decay exponent w_{20} (the w_i denote the elements, 0-indexed, of a fitted 21-parameter weight vector $\mathbf{w}$):

$$R(t, S) = \left(1 + \text{factor} \cdot \frac{t}{S}\right)^{-w_{20}}, \quad \text{factor} = 0.9^{-1/w_{20}} - 1, \quad (\text{B.3})$$

where the factor term is chosen precisely so that $R(S, S) = 0.9$ always holds, preserving stability's definition regardless of the fitted value of w_{20} . This is itself the product of two earlier simplifications: the original FSRS v4 fixed the exponent at $w_{20} = 1$ with factor = 1/9, reducing Equation B.3 to $R(t, S) = (1 + t/9S)^{-1}$; FSRS-4.5 refined this to a fixed exponent of 0.5; only from FSRS-5 onward is the decay shape itself learned from data.

Initial state and difficulty updates. A word's very first review initialises stability and difficulty directly from the grade it receives:

$$S_0(G) = w_{G-1}, \quad D_0(G) = w_4 - e^{w_5(G-1)} + 1, \quad (\text{B.4})$$

so that, for instance, an item first rated *again* ($G = 1$) starts at the lowest stability w_0 and the highest difficulty $D_0(1) = w_4$. On every subsequent review, difficulty drifts by a grade-dependent increment and is then pulled back towards a fixed target via *mean reversion*:

$$D' = D - w_6(G-3) \cdot \frac{10-D}{9}, \quad D'' = w_7 D_0(4) + (1-w_7) D', \quad (\text{B.5})$$

anchoring difficulty to $D_0(4)$ (the initial difficulty of an *easy*-rated item) rather than letting it drift unboundedly after a long run of easy or hard grades — informally termed avoiding “ease hell” in the flashcard-scheduling community.

Stability updates. Following a successful review ($G \in \{2, 3, 4\}$), stability increases according to

$$S'_r(D, S, R, G) = S \cdot \left(e^{w_8(11-D)} S^{-w_9} (e^{w_{10}(1-R)} - 1) \cdot w_{15}^{[G=2]} \cdot w_{16}^{[G=4]} + 1 \right), \quad (\text{B.6})$$

while a lapse ($G = 1$) instead resets stability to a smaller post-lapse value,

$$S'_f(D, S, R) = w_{11} D^{-w_{12}} ((S+1)^{w_{13}} - 1) e^{w_{14}(1-R)}. \quad (\text{B.7})$$

Writing $S_{\text{inc}} = S'_r/S$ for the multiplicative stability gain of a successful review, three qualitative properties of Equation B.6 recur throughout the FSRS documentation and are directly interpretable in terms of the theory already established in this chapter: S_{inc} decreases as difficulty D increases (harder items consolidate more slowly); S_{inc} decreases as stability S itself increases (an already well-consolidated memory is harder to strengthen further, consistent with the diminishing-returns shape of the Ebbinghaus curve in Section B.2.4.1); and S_{inc} increases as retrievability R decreases at the time of review, i.e. a successful recall of an *almost-forgotten* item strengthens memory more than reviewing an item that was never at risk — the computational signature of the spacing effect.

Same-day reviews (FSRS-6). The formulas above govern reviews spaced at least a day apart. FSRS-6's specific contribution is a separate, more conservative update for a card reviewed again on the same day it was last seen:

$$S'(S, G) = S \cdot e^{w_{17}(G-3+w_{18})S^{-w_{19}}}, \quad (\text{B.8})$$

which grows quickly while S is still small and flattens as S grows, converging once $S_{\text{inc}} = e^{w_{17}(G-3+w_{18})S^{-w_{19}}} = 1$; earlier versions used the same functional form without the $S^{-w_{19}}$ damping term, allowing same-day repetition to inflate stability unrealistically.

This finer-grained decomposition is, at the time of writing, the most empirically accurate publicly documented model of spaced-repetition scheduling. Critically for a new application, all of Equations B.3–B.8 update the per-item state *incrementally, after each individual review*, using a fixed parametric rule rather than a regression re-fitted in a single offline pass over a large historical corpus, as HLR requires (Section B.2.4.3). A new item begins from the initial-state Equation B.4 and refines its own difficulty and stability purely from the learner's own subsequent review history, without needing Duolingo-scale historical data first; *Glosy* adopts the published default parameter vector $\mathbf{w}$ rather than fitting its own, precisely because it does not yet have the review corpus such a fit would require. This is the specific reason Chapter C, Section C.4 (Stage 2) adopts the FSRS formulation, rather than HLR, as the basis of the `SRS_Urgency` scoring term in the recommendation engine: a word's recommendation urgency rises as its predicted retrievability R (Equation B.3) falls.

B.2.5 Active Recall and the Testing Effect

A separate, complementary strand of cognitive psychology addresses not *when* to review an item, but *how*. The testing effect — robustly demonstrated across decades of memory research — shows

that the act of actively retrieving an item from memory (as in a test or a recall task) produces substantially better long-term retention than an equivalent amount of time spent passively re-studying the same item [42]. This finding motivates a specific design constraint distinct from, and additional to, the scheduling question addressed by Section B.2.4: *when* an item is due for review, the review task itself should require productive retrieval (e.g., recalling and typing a word) rather than passive recognition (e.g., selecting a familiar-looking option among distractors). Section B.3.2 reviews game mechanics through this lens.

B.3 Modalities for Language Acquisition

Sections B.2.2–B.2.5 establish *what* an effective input and reinforcement episode should look like theoretically: comprehensible ($i+1$), authentic, cognitively bounded, spaced, and actively recalled. This section reviews the empirical literature on the two content modalities through which *Glosy* delivers input and reinforcement respectively: short-form video, addressed to the input side of acquisition, and serious games, addressed to the active-recall side.

B.3.1 Short-Form Video in Second Language Acquisition (SLA)

B.3.1.1 Incidental Vocabulary Acquisition through Audiovisual Input

A substantial body of research in Computer-Assisted Language Learning (CALL) predates the short-form video format itself and concerns the closely related case of captioned television and film viewing. Vanderplank’s early work on teletext subtitles found that captioned video supports vocabulary and listening development beyond what audio alone provides, by allowing the learner to cross-reference the written and spoken forms of an unfamiliar item in real time [43]. More recent corpus-informed work confirms this incidental-learning effect at scale: Peters and Webb found measurable incidental vocabulary gains from watching L2 television programmes without any explicit study task, with gains moderated by frequency of item occurrence and by the presence of captions [44], while Montero Perez, Peters, and Desmet found that enhancing target vocabulary within video captions (e.g., through highlighting) further improves retention relative to unenhanced captioned video [45]. Taken together, this literature establishes that passive audiovisual viewing is a genuine, if modest, channel of incidental vocabulary acquisition, and that the effect is strengthened, not weakened, by textual scaffolding synchronised to the audio track — precisely the

token-level, timestamped transcript structure that a *Glosy* reel provides (Chapter C, Section C.3).

B.3.1.2 The Gap: Passive Consumption without Reinforcement

The incidental-acquisition literature reviewed above is consistent in one further respect: the vocabulary gains it reports from passive viewing alone are real but small, typically requiring many repeated incidental encounters of an item before it is durably retained [44]. Passive viewing, in other words, is an effective delivery mechanism for comprehensible input (Section B.2.2) but is not, by itself, a spaced-repetition system in the sense of Section B.2.4: a viewer has no mechanism by which an encountered item is tracked, scheduled for review, or actively retrieved once the video ends. This is the specific gap in generic short-form social video consumption — as opposed to a purpose-built application — that motivates the research question posed in Chapter A: watching unstructured reels on an existing social platform provides input without a tracked vocabulary state, and therefore without the scheduled, active-recall reinforcement that Sections B.2.4 and B.2.5 identify as necessary for durable retention. Closing this gap requires attaching a conventional SRS review step — of the kind ordinarily triggered by a static flashcard in an application such as Anki — directly to the moment a due item is encountered inside a reel, rather than leaving that encounter to pass by unrecorded. Chapter C details the interface through which *Glosy* does this: a due word appearing in a reel’s transcript prompts an in-context review step immediately after viewing, through which the learner’s recall grade G is captured and fed into the FSRS update rules of Section B.2.4.4 (Equations B.6–B.7) — the reel itself standing in for the static flashcard as the review’s triggering stimulus.

B.3.2 Serious Games and Gamification in Language Learning

B.3.2.1 Game-Based Learning versus Gamification

Section B.1.2 (item 4) distinguished *gamification* — decorating an existing task with game-design elements such as points and streaks [28] — from *game-based learning*, in which the instructional content is embedded within the mechanical loop of an actual game [29]. This distinction is well established as a specific research strand within CALL: Godwin-Jones’s review of games in language learning documents a persistent gap between the two, noting that commercially successful language-learning products have historically favoured drill-based gamification because it is comparatively cheap to build and easy to monetise, while games in which linguistic competence is

itself the win condition remain comparatively rare [46]. Sykes and Reinhardt’s broader treatment of digital games in second and foreign language teaching similarly argues that a game’s pedagogical value derives from whether target-language production or comprehension is load-bearing to *winning the game*, not merely rewarded upon completion of an unrelated exercise [47].

B.3.2.2 Difficulty Adaptation in Hypercasual Games

Section A.2.2 (Chapter A) established the hypercasual genre’s core appeal: minimal onboarding, immediately intuitive mechanics, and short session length — properties that align closely with the microlearning argument of Section B.2.3. Within game design more specifically, dynamic difficulty adjustment (DDA) — algorithmically tuning a hypercasual game’s difficulty in response to the individual player’s observed performance, rather than fixing it in advance — has been shown to improve player retention and perceived replay value relative to a static difficulty curve [48]. This provides the game-design-theoretic counterpart to the lexical personalisation argument of Section B.4: just as a reel should be selected for lexical comprehensibility, a game level’s difficulty should be tuned to the player’s demonstrated skill — in *Glosy*’s case, the specific vocabulary items the level is built from, rather than a generic difficulty parameter.

B.3.2.3 Active-Recall Game Mechanics for Vocabulary

The testing-effect literature reviewed in Section B.2.5 [42] favours game mechanics that require *productive* retrieval of a target word’s orthographic form over mechanics that require only its *recognition* among distractors. Word-reconstruction puzzle formats — in which the player must recall and produce a target word letter-by-letter, validating each attempt against positional or compositional constraints — structurally enforce this kind of active retrieval as the core win condition of the game itself, rather than as an optional or bypassable step. This is the specific theoretical property, rather than genre popularity alone, that motivates the word-reconstruction game modalities implemented in *Glosy* and documented in Chapter C, Section C.4.

B.4 Algorithmic Frameworks for Personalized Recommendation

Sections B.2 and B.3 establish *which* content is pedagogically appropriate for a learner at a given moment: comprehensible (Section B.2.2), due for review (Section B.2.4), and appropriately re-

inforcing (Section B.3). Selecting such content automatically, per learner, is a recommendation problem in the information-retrieval sense. This section reviews the two classical recommendation paradigms, content-based and collaborative filtering, and the strategies the literature offers for combining them; Section B.5 maps this theory onto *Glosy*'s own recommendation engine.

B.4.1 Content-Based Filtering

Content-based filtering recommends items by comparing an item's own features against a profile of a single user's own preferences, built from that user's history, without reference to other users' behaviour [49]. Its principal known limitation is that it has no way to express a temporal signal such as *when* a previously-seen item is due for review [49] — a gap the memory-decay models of Section B.2.4 are specifically designed to fill.

B.4.2 Collaborative Filtering and Hybrid Recommender Systems

Collaborative filtering (CF) recommends items based on patterns of agreement across many users' behaviour — favouring items that users with a similar engagement history have responded well to — rather than on an item's own features [49]. Because it depends on an interaction matrix populated across many users, CF is known to perform poorly under the *cold-start* problem, when a system has few users or little accumulated interaction history [49].

Burke's survey of hybridisation strategies for combining multiple recommendation techniques distinguishes several ways of doing so, including *weighted* hybridisation, in which each technique produces an independent score and the final score is their linear combination, and *cascade* hybridisation, in which one technique coarsely ranks or partitions the candidate set and a second technique is applied only to break ties within it [50]. Ricci et al.'s handbook situates both strategies, and content-based and collaborative filtering more generally, within the broader design space of recommender systems [51].

B.4.3 Recommender Systems in Educational Technology

Recommender systems applied to e-learning typically personalise the sequencing of resources or exercise *difficulty* to a learner's inferred ability, rather than the semantic content of what is recommended [52] — the same gap identified in Section B.1.2 (item 1).

B.5 Synthesis: A Theory-Driven Design Framework

This chapter has reviewed, in turn, the pedagogical theory of comprehension and retention (Section B.2), the empirical literature on the two content modalities through which *Glosy* delivers input and reinforcement (Section B.3), and the algorithmic frameworks by which content is selected and sequenced for the individual learner (Section B.4). Table B.2 makes explicit the traceability between each theoretical foundation reviewed above and the corresponding design decision it motivates in Chapter C, operationalising the theory-driven design approach stated at the outset of Section B.2.

Table B.2: *Theory-driven design traceability matrix: theoretical foundation, design implication, and the corresponding Glosy component.*

Theoretical Foundation	Design Implication	Glosy Component
ZPD (B.2.1) & Input Hypothesis (B.2.2)	Recommend content at the $i + 1$ comprehensibility boundary	Stage 1: Comprehensibility Filtering, 98% lexical coverage (Sec. C.4.1)
Cognitive Load Theory & Microlearning (B.2.3)	Bound each learning episode to a short, self-contained unit	90-second reel cap; Wordle attempt limit; Word of Wonders word/letter caps (Sec. C.1, Sec. C.6)
Forgetting Curve & Spaced Repetition (B.2.4)	Schedule review before predicted recall failure	Stage 2: Spaced-Repetition Prioritisation, due-word partitioning (Sec. C.4.2)
Testing Effect (B.2.5) & Game-Based Learning (B.3.2)	Require productive retrieval, not passive recognition	Wordle-style and Word of Wonders vocabulary games (Sec. C.6)
Cascade Hybrid Recommender (B.4.2)	Filter by comprehensibility, partition by due-for-review, break ties by engagement-weighted content similarity	Stage 1–3 cascade ranking (Sec. C.4.4)
Sociocultural Theory (B.2.1)	Provide social scaffolding for culturally-opaque authentic input	Collaborative learning features (Ch. F, Future Work)

Read together with the platform limitations identified in Section B.1.2, this synthesis reframes the research questions posed in Chapter A as a design problem with a specific theoretical answer for each of its parts: RQ1 is addressed by the content-based lexical filtering of Section B.4.1 operating

over the short-form video modality of Section B.3.1; RQ2 is addressed by the cascade hybrid of Section B.4.2, fusing comprehensibility with the memory-decay models of Section B.2.4; and RQ3 is addressed by the active-recall game theory of Section B.3.2, parameterised by the same memory-decay state that governs RQ2. Chapter C now presents the system architecture that implements this framework.

Chapter C

System Architecture & Methodology

C.1 System Overview

Glosy is designed as a three-service architecture in which a mobile frontend, a Node.js backend API, and a dedicated Python reels microservice each own a distinct part of the system's responsibilities, all three sharing a single PostgreSQL database. In development, the three services and the database are orchestrated together with Docker Compose: the backend API listens on port 3500, the reels service on port 3600, and PostgreSQL 16 on its default port, each built and reloaded independently of the others so a change to one service does not require restarting the rest of the stack.

- **Frontend:** A cross-platform mobile application built with React Native and Expo Router, targeting Android and iOS from a single JavaScript/TypeScript codebase, and responsible for the offline-first client state, the reels feed, and the two vocabulary-driven games.
- **Backend API:** A Node.js/Express service responsible for authentication, user and vocabulary state, offline synchronisation, and the dictionary and letters modules.
- **Reels Service:** A Python/FastAPI microservice responsible for serving reels together with their structured dialogue data, and for the three-stage recommendation engine that selects and ranks them per learner.

- **Database:** A single PostgreSQL 16 instance, shared by both back-end services, whose schema encodes linguistic content, user state, and engagement signals (Section C.3).

Splitting the reels service from the backend API is a deliberate choice rather than an accident of the codebase’s history: the recommendation engine’s queries — joining a reel with its dialogue, sentences, per-token word data, and translations, then scoring the result against a learner’s vocabulary and engagement history — are naturally expressed as an object-relational query built up in code, which SQLAlchemy’s async ORM supports directly, whereas the backend API’s own queries are simple enough that raw parameterised SQL via the `pg` client remains the more direct choice. The two services therefore use different data-access strategies deliberately matched to the shape of the queries each one runs, rather than a single ORM standardised across the whole system.

A reel’s duration is capped at 90 seconds at upload time, rejected outright by the reel-creation schema if exceeded; this is a direct application of the microlearning bound of Section B.2.3, keeping a single unit of video input short enough that it cannot itself exceed a learner’s working-memory budget regardless of how much unfamiliar vocabulary it contains. Section C.6 describes the analogous bound on the other side of a learning episode, the games’ own attempt and word-count limits.




Figure C.3: *High-level system architecture: the mobile frontend communicates with the Backend API (`api.glosy.gr`) and the Reels Service (`recs.glosy.gr`) over authenticated REST endpoints; both services share a single PostgreSQL instance.*

C.2 Technology Stack

C.2.1 Frontend — React Native / Expo

The frontend is built with React Native 0.86 and React 19, on Expo SDK 57, using Expo Router for file-system-based navigation between the application’s five tabs and its onboarding, settings, and

game routes. `expo-video` handles reel playback, `expo-secure-store` holds authentication tokens, and React Native Reanimated (4.5) drives the gesture-based interactions in the reels feed and the two games, together with `@gorhom/bottom-sheet` for the swipe-to-review prompt described in Section C.5.2. The FSRS scheduling itself (Section C.4.4) runs entirely on-device via `ts-fsrs`, a published TypeScript implementation of the algorithm reviewed in Section B.2.4.4, rather than a bespoke implementation — there is no pedagogical reason to re-derive difficulty/stability/retrievability update rules whose reference implementation is both open-source and already exercised against the FSRS benchmark.

C.2.2 Backend API — Node.js / Express

The backend is a Node.js service on Express 4, communicating with PostgreSQL through the plain `pg` client rather than an ORM, since its own queries (per-row inserts/updates keyed by user and word id) do not benefit from an object-relational mapping layer the way the reels service’s join-heavy queries do (Section C.1). Authentication combines JSON Web Tokens (short-lived access tokens, longer-lived refresh tokens, rotated via `jsonwebtoken`) with Google OAuth 2.0 (`google-auth-library`) as an alternative sign-in path; every request body is validated against a Joi schema before it reaches a controller, and the API is self-documented with Swagger/OpenAPI at `/swagger`. Video thumbnails are generated server-side with `fluent-ffmpeg` over a statically-bundled `ffmpeg` binary, and both the backend and the reels service read and write the same Cloudflare R2 (S3-compatible) object storage bucket through the AWS SDK’s `S3Client`, so a reel’s media is addressable identically regardless of which service produced it.

C.2.3 Reels Service — Python / FastAPI

The reels service is a FastAPI application over Pydantic 2 for request/response validation and SQLAlchemy 2’s async ORM (via `asyncpg`) for the join-heavy queries described in Section C.1. It does not run its own authentication: instead, `python-jose` decodes the same HS256-signed access tokens the Node backend issues, against the same shared secret, so a single login session is valid across both services without a second sign-in step or a token-exchange call between them. Stage 3 of the recommendation engine (Section C.4.3) is implemented with scikit-learn’s `TfidfTransformer` and `cosine_similarity`, a deliberately lightweight, CPU-only choice: the service runs on a 512 MB host, which rules out loading a full NLP pipeline (spaCy, Stanza, or Hazm) into every request worker purely to vectorise a reel’s word-id bag. Sentence to-

kenisation and lemmatisation, which do need those heavier libraries, are instead performed offline by a separate backfill script rather than at request time, keeping the request path itself dependent only on scikit-learn’s lightweight vectorisers.

C.2.4 Database — PostgreSQL

Both backend services share a single PostgreSQL 16 instance (Section C.1), chosen for its support of array operators, window functions, and common table expressions — used throughout the vocabulary bulk-seeding logic of Section D.2.2 — over a lighter-weight embedded database that would not support these directly in SQL. The full schema is detailed in Section C.3 and Section D.1.

C.3 Data Modelling

The schema is organised along the same three axes used throughout Chapter B’s theoretical framework: content that exists independently of any learner, the state the system tracks about a specific learner, and the signals a learner generates while engaging with that content.

C.3.1 Linguistic Content

This axis holds everything a learner can be shown, independent of who they are: `languages`, `words` (each tagged with a `level` drawn from the CEFR scale A1–C2 plus a small number of non-CEFR codes used elsewhere in the schema) and their `word_translations`, `letters` for the alphabet-drill dictionary entries, `sentences` with per-token (`sentence_tokens`) and per-language (`sentence_translations`) annotation, `dialogues` that group sentences into a single reel’s transcript, and `reels` themselves, which reference a dialogue and their creator but otherwise carry only media metadata (URL, thumbnail, duration). A reel’s own row never records who it is comprehensible *to*: comprehensibility is only ever computed relative to a specific learner’s vocabulary at query time (Section C.4.1), not stored as a property of the reel itself, even though the reel row does reference the single user who created it.




Figure C.4: Entity-relationship diagram of the 21-table schema, grouped into the three axes of Section C.3: Linguistic Content, User State, and Engagement Signals.

C.3.2 User State

`users` carries authentication and profile fields alongside the two progression currencies from Section C.5.1, `coins` and `energy`. `user_languages` records one row per (user, native language, learning language) triple, including a `proficiency_level`, a per-language `experience` counter, and an `is_current_language` flag, since a learner may study more than one language but only one is active at a time. `user_vocabulary` is the learner model proper: one row per (user, word) pair, carrying the full FSRS memory state reviewed in Section B.2.4.4 directly as columns — `stability`, `difficulty`, `lapses`, `fsrs_state`, `learning_steps`, `review_count`, and `next_review_at` — so that “is this word due” is a plain column comparison rather than a value recomputed from a review history log on every read.

C.3.3 Engagement Signals

`reel_interactions` records one row per (user, reel) pair, accumulating likes, saves, shares, and comments as boolean/text columns rather than a separate event-log table, plus the two book-keeping fields Stage 1 and Stage 3 depend on directly: a `view_count` distinguishing a genuine repeat view from the first, and a `last_recommended_at` timestamp used to exclude a reel already shown to this learner within the cooldown window (Section C.4.1). `reel_reports` and `feedback` complete this axis, capturing moderation and product-feedback signals that do not feed the recommendation engine but share the same “signal generated by a learner about a piece of content” shape.

C.4 Recommendation Engine Design

The engine’s goal is to select, from the pool of reels in the learner’s target language, those that are simultaneously comprehensible, due for review, and matched to that learner’s own demonstrated interests. It does so as a three-stage pipeline over a candidate pool scoped to the learner’s language pair: a hard filter, followed by two re-ranking passes over the filter’s survivors.

C.4.1 Stage 1: Comprehensibility Filtering

Every candidate reel is first checked against two independent conditions, either of which removes it from consideration outright. The first is recency: a reel already shown to this learner within a

cooldown window is excluded, tracked via a `last_recommended_at` timestamp rather than an actual-viewing signal, since a genuine view is only confirmed after 2 seconds of continuous playback and would otherwise let an unwatched, scrolled-past reel block itself from ever being shown again. The second is the comprehensibility threshold operationalising Section B.2.2’s $i + 1$ boundary directly:

$$\text{Coverage}(U, R) = \frac{|V_R \cap V_U|}{|V_R|} \geq 0.98 \quad (\text{C.1})$$

where V_R is the unique set of word ids in reel R ’s transcript and V_U is the set of word ids anywhere in learner U ’s tracked vocabulary — membership only, independent of each word’s FSRS state, since comprehension depends on a word being recognisable at all, not on how due it is for review. This is the content-based filtering pattern reviewed in Section B.4.1: a reel’s own lexical features scored against a profile built from this one learner’s vocabulary, with no reference to any other learner. Every reel surviving both conditions carries its computed coverage forward as a `comprehensibility_percentage`, displayed to the learner alongside the reel itself.

C.4.2 Stage 2: Spaced-Repetition Prioritisation

Because FSRS runs on-device (Section C.2.1), the recommendation engine does not compute due words itself: the frontend supplies an ordered list of word ids currently due for review, oldest-due-first, with each request. Among Stage 1’s survivors, a reel containing at least one of these words is preferred over one that does not — re-ranking the survivor set rather than filtering it further, so that a learner is never shown fewer reels merely because none of them happen to contain a due word. This is the mechanism that closes the gap identified in Section B.3.1.2: watching a reel that happens to contain a due word is what triggers the in-context review step of Section C.5.2, giving passive viewing the scheduled, active-recall reinforcement it otherwise lacks. Once a page of reels is selected, due words are assigned back to specific reels on a first-due, first-assigned basis, so that a page carrying several due words can surface a distinct review prompt for each, rather than only ever reviewing the single most-overdue word.

C.4.3 Stage 3: Content-Based Personalisation

Within equal Stage 2 priority, survivors are further ranked by similarity to the learner’s own engagement history — a second content-based pass, not collaborative filtering (Section B.4.2), since no other learner’s data enters the computation. Each reel in the candidate pool is represented as

a TF-IDF vector over its word-id bag (Section C.2.3 motivates the choice of scikit-learn here on engineering, not only theoretical, grounds); the learner’s profile vector is the weighted *sum* (not an average — there is no normalisation by total weight or interaction count) of the vectors of reels they have previously engaged with, weighted by the strength of that engagement (a comment or a save contributing more than a bare view), and candidates are ranked by cosine similarity to that profile. A learner with no engagement history yet simply has no profile vector, and this stage contributes no signal for them, rather than failing or substituting an arbitrary default.

C.4.4 Combining the Three Stages: A Cascade, Not a Weighted Sum

The three stages are combined as a *cascade* in Burke’s sense (Section B.4.2), not a weighted linear combination: Stage 1 is a hard pass/fail filter, and Stages 2 and 3 together form a single sort key, (has-due-word, content-similarity), applied to Stage 1’s survivors with a random tiebreak. Stage 2’s binary flag is always the primary key and Stage 3’s similarity only breaks ties within it, so a reel containing a due word is never outranked by one that merely resembles the learner’s past engagement more closely. Both signals degrade gracefully to a constant when they have nothing to work with — no word currently due, or no engagement history yet to build a profile from — which turns that half of the sort key into a no-op rather than a failure; when neither stage has any signal at all, the ranking degrades cleanly to the same random selection a brand-new learner would have received before either signal existed.

C.5 Gamification Design

Section B.3.2.1 distinguished decorative gamification — points and streaks layered onto an unchanged drill — from game-based learning, in which the reward layer is structural. *Glosy*’s progression system is deliberately scoped to the former role only: a light extrinsic-motivation layer sitting on top of the two games and the reel-review loop, which are themselves where the game-based-learning argument of Chapter B is made.

C.5.1 Progression Currencies: Coins, Energy, and Experience

Three progression stats are tracked per user (or, for *experience*, per learning language): coins, energy, and experience. Of the three, only coins currently have a complete reward loop: a learner

earns coins for rating a spaced-repetition review prompt after watching a reel (Section C.5.2) and for winning a round of either game. The coin amounts are not yet unified into a single reward schedule, however — the review prompt and a Word of Wonders win both award a flat 20 coins, while Wordle awards 5 coins for a win and 1 for a loss — a discrepancy worth reconciling into a single schedule rather than a deliberate design choice.

Energy and experience, by contrast, are present in the data model, synced between client and server, and displayed in the profile and header UI, but have no gameplay loop attached to them yet: no screen in the application currently spends energy to start a round, and no action currently increments experience. The schema and configuration already anticipate more than is implemented — `users` carries an `energy_updated_at` timestamp alongside `coins_updated_at`, and the backend defines `MAX_ENERGY` and `ENERGY_REGEN_INTERVAL` environment variables — but no code path on either the client or the server currently reads them to compute a regenerated value, and the stored default for a new user’s energy (100) does not match the validated cap those settings imply (5). These two currencies should be read as scaffolding for a future energy-gated play loop, not as an oversight discovered late; Chapter F returns to this as a direction for future work.

C.5.2 Design Principle: Every Game Round Doubles as an FSRS Review

Both vocabulary-driven games and the reel-review prompt share one design rule: the word being tested is always drawn from the learner’s own tracked vocabulary, and preferentially from the words FSRS currently marks as due, rather than from the full dictionary or an arbitrary in-game selection. This is the direct operationalisation of Section B.2.5’s testing effect and the gap identified in Section B.3.1.2: a game round is not merely themed around vocabulary, it *is* the productive-retrieval review step that closes that gap, with the game’s win condition and the FSRS review coinciding rather than sitting side by side.

The two games apply this rule differently, however, rather than sharing one uniform mechanism. Word of Wonders only ever treats the practice word as a graded FSRS review, surfacing the learner-facing Forgot/Hard/Good/Easy prompt of Section C.5.2 on its finish screen, when that word actually came from the due-word queue — finding the board’s practice word when nothing was due is not graded at all. The Wordle-style game does not use that self-graded prompt in any case: finding the secret word within the attempt limit is always recorded as a successful review (a fixed **GOOD** rating) regardless of whether the word was genuinely due or a random tracked fallback, and running out of attempts leaves the word’s FSRS state untouched rather than recording a failed review.

Word of Wonders' gate exists precisely to avoid the corruption a uniform rule would otherwise risk — grading a word that was never due would inject a review event that should not exist — while Wordle accepts that risk in exchange for a simpler, ungraded win/loss outcome; Section D.5.1 and Section D.5.2 detail each mechanism concretely.

C.6 Vocabulary-Driven Games

Both games realise the word-reconstruction format motivated in Section B.3.2.3: the player recalls and produces a word letter-by-letter rather than recognising it among distractors, and in both games that word is drawn from the learner's own tracked vocabulary under the rule of Section C.5.2. Each is also bounded in scope for the same microlearning reason as the 90-second reel cap of Section C.1: Wordle limits a round to six guesses, and Word of Wonders caps both the number of words placed on a board and the size of the shared letter pool a learner must scan (Section D.5.2), so neither game's single round can grow into a long-form session at odds with Section B.2.3's account of bounded working memory.

C.6.1 Wordle-Style Game

The learner has six attempts to guess a five-letter secret word, with each guess coloured letter-by-letter as correct, present-elsewhere, or absent, following the now-familiar Wordle mechanic. What distinguishes it from a generic port of that mechanic is that it is genuinely multilingual rather than English-only with translated labels: each of the three supported languages (English, Greek, Farsi) has its own letter-validity pattern restricting guesses to that language's script, its own on-screen keyboard layout, and, for Farsi, a right-to-left board and keyboard rather than a mirrored English one. This matters pedagogically as well as cosmetically: a Greek or Farsi learner is guessing within their target language's own orthography, not a Latin-script approximation of it.

C.6.2 Word of Wonders (Crossword)

Word of Wonders places the practice word — the due word, or a random tracked fallback (Section C.5.2) — on a grid alongside several other words from the same learner's tracked vocabulary, each intersecting the others at a shared letter, in the style of a crossword. Where Wordle exercises one word per round, Word of Wonders' board therefore reinforces several previously-introduced

words simultaneously around the single word under FSRS review, at the cost of a puzzle that takes visibly longer to solve; the concrete placement algorithm, and the letter-budget constraint that keeps that trade-off bounded, are detailed in Section D.5.2. A secondary “extra words” mechanic rewards the learner for continuing to explore the revealed letters after the placed words are found — any additional valid word spelt out by the board’s letters counts towards a bonus, paid out in batches — extending engagement with the same board past the point where the due-word review is technically already complete.

Chapter D

Implementation Details

D.1 Database Schema

The schema comprises 21 tables in a single PostgreSQL 16 database (Section C.2.4), organised below along the same three axes introduced in Section C.3.

D.1.1 Linguistic Content Tables

`languages`, `words` and `word_translations`, and `letters` hold the base vocabulary. A sentence’s text lives in `sentences`, its per-word annotation in `sentence_tokens` (word id, position, part of speech), and its cross-language equivalents in `sentence_translations`. `dialogues` groups an ordered sequence of sentences via `dialogue_sentences` (which also carries each sentence’s `start_time_ms/end_time_ms` within the video), and `reels` references exactly one dialogue alongside its media URL, thumbnail, duration, and creator. Rather than re-joining `sentences`, `sentence_tokens`, and `sentence_translations` on every read, `dialogues` additionally carries a denormalised `sentences_json` column: a pre-computed snapshot of the full dialogue, tokens and all-language translations included, written once when a reel’s subtitles are created or edited and read directly thereafter (Section D.4.2).

D.1.2 User State Tables

`users` holds authentication fields and the two implemented-and-scaffolded progression currencies of Section C.5.1 (`coins`, `energy`, each with its own `_updated_at` timestamp). `user_languages` records one row per language pair a user studies, with its own `proficiency_level` and `experience`. `user_vocabulary` is the FSRS learner model described in Section C.3.2, carrying `stability`, `difficulty`, `lapses`, `fsrs_state`, and `learning_steps` directly as columns. `user_sentences` tracks per-sentence progress on a separate, coarser scheme: a `mastery_level` between 1 and 6, alongside its own `review_count` and `next_review_at`, distinct from `user_vocabulary`'s continuous FSRS state. `refresh_tokens` (family-scoped, for the rotation scheme of Section D.2.1), `password_reset_codes`, and `email_verification_codes` round out this axis with authentication-adjacent state.

D.1.3 Engagement Signal Tables

`reel_interactions` carries one row per (user, reel) pair: `boolean` `is_liked/is_saved/is_shared` flags, an optional comment, a `view_count`, and the `last_recommended_at` timestamp Stage 1 filters on (Section C.4.1). `reel_reports` and `feedback` record moderation reports and free-form product feedback respectively; neither feeds the recommendation engine, but both share this axis's shape of a learner-generated signal attached to a specific piece of content.

D.2 Backend API

The API is documented with Swagger/OpenAPI, generated from JSDoc annotations on each route and served at `/swagger`. Every request body is validated against a Joi schema before it reaches a controller (Section C.2.2), and large, uniform list responses — the dictionary word list and a learner's vocabulary state — are serialised in a columnar `{columns, rows}` shape rather than one JSON object per row, avoiding repeated field names across thousands of otherwise identically-shaped entries.

D.2.1 Authentication and Authorization

D.2.1.1 JWT Access/Refresh Token Rotation

Login and registration issue a short-lived access token and a longer-lived refresh token. Rather than storing refresh tokens verbatim, the `refresh_tokens` table stores only a hash of each one, grouped into rotation *families* by a `family_id`. Every use of `POST /refresh` runs a single atomic operation: it locks the presented token's row (`SELECT . . . FOR UPDATE`, inside a transaction, so two concurrent refresh attempts against the same token cannot race each other), and then branches on that row's state:

- if the row is already `revoked`, the presented token has already been rotated away once before — a replay of an old token — so every active token in that family is revoked immediately and the request is rejected;
- if the row is unrevoked but past `expires_at`, the family is revoked and the request is rejected;
- otherwise, a new token row is inserted into the same family, copying the original `expires_at` verbatim so a session's absolute lifetime never resets on activity, the presented row is marked revoked with a `replaced_by_id` pointer to the new one, and the new access/refresh pair is returned.

This gives the family reuse-detection property standard to refresh-token rotation: a stolen refresh token is only usable once before the legitimate client's next refresh attempt discovers the token has already moved on and revokes the whole family, rather than the theft going undetected indefinitely.

D.2.1.2 Google OAuth 2.0

`verifyGoogleToken` middleware validates a client-supplied Google ID token against the OAuth client id registered for the requesting platform, using `google-auth-library's OAuth2Client.verifyIdToken`, before either the login or registration controller runs. On success it attaches the decoded profile (Google id, email, name, picture) to the request; the downstream controller then finds-or-creates a user keyed by that Google id exactly as it would for a password-based account, so both sign-in paths converge on the same JWT issuance logic described above.

D.2.2 Registration and Vocabulary Initialisation

Registration requires a verified email address: a code sent to the address must be confirmed against a hashed, expiring, attempt-limited pending record before any user row is created, and the pending record is deleted on success so the code cannot be reused. The username is derived from the email's local part, with two disambiguating fallback variations (an underscore or dot plus a numeric suffix) tried in sequence if the plain form is already taken.

D.2.2.1 Proficiency-Based Vocabulary Auto-Seeding

Immediately after a new language pair is created, if the learner declared a proficiency level above the lowest, `userVocabularyModel.addByProficiencyLevel` seeds `user_vocabulary` with every word at a level below their stated one, on the assumption that a self-declared B1 learner already knows their A1 and A2 vocabulary without having reviewed it inside *Glosy*. Seeded words are inserted directly into FSRS's Review state rather than New, with an initial stability drawn — uniformly at random per word, via a single SQL CASE expression keyed on each word's own level — from a range that widens with distance below the target level (up to 7 days for the level immediately below, 20 for two levels below, 60 for three or more): a word closer to the learner's stated level is assumed more recently and less thoroughly reinforced than one several levels more elementary. Randomising per word rather than assigning one flat value per level matters operationally as much as pedagogically: without it, an entire level's worth of seeded words — potentially thousands, for a learner who declares a high starting level — would all share the same `next_review_at` and become due for review on the exact same day.

D.2.3 Offline Synchronisation

A single `POST /sync` endpoint reconciles everything the frontend's offline-first client (Section D.3.3) has accumulated since its last successful sync: updated `user_progress` (coins, energy), and `vocabulary_changes/sentence_changes`, each split into inserts, updates, and deletes keyed by word or sentence id.

D.2.3.1 Change-Set Ordering: Deletes, Updates, Inserts

Within each change-set, the three operations are applied in a fixed order — deletes, then updates, then inserts — rather than in whatever order the client happened to record them. This ordering is what makes the endpoint correct in the presence of a delete-then-recreate on the same word or sentence id during one offline session (a learner removing a tracked word and later re-adding it before ever reconnecting): applying the delete first and the insert last guarantees the row exists with the client’s latest state once the request completes, regardless of the order the two actions occurred in on-device. Vocabulary updates are currently applied one word at a time in a per-row loop rather than as a single batched statement; Chapter F identifies replacing this with a bulk `INSERT . . . ON CONFLICT DO UPDATE` as a concrete performance target for large offline sessions.

D.2.4 Dictionary and Letters Modules

The dictionary endpoints serve a language’s full word list, optionally joined with translations into a second language, keyed by ISO language code and validated against a fixed three-language enum (English, Farsi, Greek) hard-coded into the request schema, rather than a live lookup against whichever languages happen to be present in the database at the time. The letters module serves each language’s alphabet as its own lightweight dictionary — vowels and consonants for Farsi split further into short/long forms, per the RTL and script-specific handling already noted for the Wordle-style game (Section C.6.1) — so that alphabet drills can reuse the same word-list UI components as ordinary vocabulary instead of a bespoke screen.

D.3 Frontend Architecture

D.3.1 Navigation and Screen Structure

Expo Router maps the file system directly onto the navigation tree. The five tabs from Section C.1 — Home, Reels, Create, Practice, Profile — sit under a single (`tabs`) route group, each independently stackable (Create, for instance, pushes a details and a subtitle-sync screen on top of its own tab without affecting the other four); onboarding, settings, and the two full-screen games each live under their own top-level route group outside the tab bar, since none of them should show tab chrome while active.

D.3.2 State Management: Context and Hooks

Vocabulary state is split across two parallel structures rather than one: `userVocabulary` itself (keyed by word id, holding each word's FSRS fields) and a separate `vocabularyChanges` object recording only the inserts/updates/deletes made since the last successful sync. A single dispatch function updates both together for an ordinary edit, through two separate reducers, so the change-set is always an accurate diff rather than something reconstructed by re-comparing full snapshots at sync time. Two call sites deliberately bypass the change-set half of this: `bulkAddVocabulary`, used for onboarding auto-fill, since that seed mirrors what the backend's own proficiency-based auto-seed (Section D.2.2) already creates server-side and so is not itself a change to push back; and `reviewWord`, which computes a word's next FSRS state on-device via `scheduleReview` (Section C.2.1) and dispatches it as an ordinary `UPDATE` — a graded spaced-repetition review is not treated as a special kind of state change requiring its own sync path, it is simply an update to the same fields any manual edit would touch.

D.3.3 Offline-First Strategy

Every piece of state that must survive an app restart — vocabulary, its pending change-set, sentences, user progress — is held through a shared `usePersistedState` hook backed by `AsyncStorage`, each exposing an `isLoading` flag so dependent logic (most importantly, sync) can wait for disk hydration to finish before deciding anything is dirty. `useBackendSync` generalises over an arbitrary list of such change-sets and pushes them to `POST /sync` (Section D.2.3) under three independent triggers: a five-minute interval while changes are pending, the instant `isOnline` transitions from false to true, and the app moving to the background — covering the periodic case, the reconnection case, and the "about to be suspended, don't lose this" case respectively, rather than relying on any one of them alone. Dirtiness itself is tracked precisely rather than optimistically: a change to `coins` or `energy` is only flagged when its actual value diverges from the last value known to be reflected server-side, so that a reference-only update to `userProgress` — switching the active language, or applying a fresh snapshot the server itself just returned — does not mark anything dirty and trigger a needless sync. A failed sync leaves its change-sets marked dirty for the next trigger to retry, surfaces an error banner once rather than on every retry, and treats an HTTP 401 as "this session is over," clearing the dirty flag rather than retrying a request that will never succeed until the learner logs in again.

D.4 Reels Service

D.4.1 Service Architecture

The reels microservice (`reels-service/`) follows a conventional FastAPI layering: `routers` parse and validate the HTTP layer only, delegating immediately to a `services` class holding the actual business logic (`ReelService` for reading/recommending reels, `ReelCreationService` for writing them), which in turn operates on SQLAlchemy `models` through Pydantic `schemas` at the boundary. `GET /reels` is deliberately dual-mode: it decodes the same JWT the Node backend issues (Section D.2.1) if one is present, but does not require it, so an unauthenticated request still gets a working reels feed via `get_random_reels` while an authenticated one gets the full recommendation pipeline via `get_personalized_reels` — there is no separate “guest mode” flag anywhere in the request, only the presence or absence of a decodable token.

D.4.2 Reel Creation and Dialogue Enrichment

Creating a reel with its subtitles is one all-or-nothing transaction: a `dialogues` row, one `sentences` row per subtitle line (reusing an existing row for identical text in that language rather than duplicating it), one `dialogue_sentences` row per line carrying its millisecond-precise start/end time, and any translation lines linked in via `sentence_translations`. A line’s position in the client-supplied array, not any position value the client sends, determines its stored `position`, so a malformed or out-of-order request can never violate the dialogue’s position ordering. Tokenising each sentence into `sentence_tokens` — which requires the lemmatisation libraries Section C.2.3 deliberately keeps out of this service’s live request path — does not happen at creation time; a newly-created reel’s dialogue therefore has no per-word token data until a separate offline script backfills it. What is computed immediately, in the same transaction, is the `sentences_json` denormalisation of Section D.1.1: one aggregating SQL query rebuilds the dialogue’s full JSON snapshot — text, every language’s translations, and whatever tokens already exist — once, at write time, rather than reassembling it from four joined tables on every subsequent read. Editing a published reel’s subtitles (`PUT /reel/{id}/dialogue`) reuses the identical insert logic over the new lines, after first clearing the dialogue’s existing `dialogue_sentences` links, and recomputes the same JSON snapshot.

D.4.3 Recommendation Engine Implementation

D.4.3.1 Stage 1: Candidate Loading and Comprehensibility Scoring

The candidate pool for a request is every reel in the learner’s target language, loaded with its language, creator, and dialogue eager-loaded in one query — the catalog is small enough at present that scoring the whole pool in application code is acceptable, a scaling assumption Section D.4.4 returns to. Two SQL-backed exclusions run first: reels already recommended to this learner within the cooldown window (a `reel_interactions.last_recommended_at` lookup), and, separately, whether the learner has a `user_languages` row for this pair at all — if not, the request falls through entirely to the same unauthenticated `get_random_reels` path a guest gets, since there is no tracked vocabulary yet to filter on. For each surviving candidate, its unique word-id set is extracted from the cached `sentences_json` in Python (no query, since it is already in memory) and intersected against the learner’s known-word set, also fetched once per request; a candidate passes only once that ratio reaches `COMPREHENSIBILITY_THRESHOLD` (0.98, Section C.4.1). A reel with no tokens yet — one whose dialogue has not been through the offline tokenisation backfill — cannot be measured against this threshold at all, and is excluded rather than assumed comprehensible.

D.4.3.2 Stage 2: Frontend-Supplied Due-Word Partitioning

The engine never computes which words are due itself: `due_word_ids` arrives as a comma-separated, oldest-first query parameter, populated from the on-device FSRS state of Section D.3.3. Stage 1’s survivors are partitioned by simple set membership — does this reel’s word-id set intersect the due-word set at all — and that boolean, not the number or identity of due words matched, is the primary key of the sort described next. Once a page is selected, due words are assigned back to specific reels oldest-first, each to the first still-unclaimed reel in that page containing it, so that a single response can carry a distinct review prompt for more than one due word rather than only ever the single most-overdue one.

D.4.3.3 Stage 3: TF-IDF Profile Vectors and Cosine Similarity

Every candidate’s word-id bag is vectorised together, over the whole pool rather than per-request, so that the TF-IDF weighting’s inverse-document-frequency term reflects the full catalog and not

just one learner’s survivors: `DictVectorizer` turns each reel’s `{word_id: count}` mapping into a sparse count matrix, and `TfidfTransformer` weights it. The learner’s profile vector is the sum of the TF-IDF rows for every reel they have previously interacted with, each row scaled by that interaction’s weight (a flat 0.2 for a bare view, stacked with 1.0 for a like, 1.5 for a save, 1.0 for a share, and 1.5 for a comment, so a single heavily-engaged reel can outweigh several passive views); a learner with no interaction history yet simply has no profile vector, and `cosine_similarity` is skipped rather than computed against a zero vector.

Stages 2 and 3 are combined as a single sort key over Stage 1’s survivors, `(has_due_word, similarity)`, applied after a random shuffle so equally-ranked candidates do not always resolve in the same order: this is Python’s ordinary stable tuple sort, not a weighted formula (Section C.4.4). When a request has neither a due-word list nor an existing profile vector to work with, both components of the key are constant across every candidate, and the stable sort degrades to the pre-shuffle random order — the same random selection `get_random_reels` would have produced.

D.4.4 Personalisation Implementation Status

Three simplifications are explicit, commented design decisions in the code rather than oversights discovered while writing this thesis. First, the candidate pool is loaded and scored in full on every request; the service’s own comment on this notes it directly: the catalog is small enough for now that this is acceptable, and it should move to a SQL-side computation once the reel catalog grows large enough for per-request, in-memory TF-IDF fitting to become a bottleneck. Second, there is no collaborative-filtering component anywhere in the engine, for the cold-start reason already given in Section B.4.2: a cross-learner interaction matrix dense enough to be useful does not yet exist for a newly-launched application, so Stage 3 substitutes a second content-based pass over the learner’s own history rather than attempting a collaborative signal on too little data. A third, narrower limitation follows directly from Section D.4.2: a reel is only ever comprehensibility-scored once its dialogue has been tokenised, so a newly-uploaded reel is effectively invisible to the personalised feed until the offline backfill script has processed it, a manual step rather than an automatic part of publishing.

D.5 Game Modules

D.5.1 Wordle-Style Game Implementation

The secret word is drawn exclusively from the learner’s own `userVocabulary` (Section C.5.2), never the full dictionary: due words are collected via `getDueWords` (Section C.2.1) and one is chosen at random when at least one exists, falling back to a random tracked word otherwise. Guesses and the secret word are compared through the same `normalizeWord` canonicalisation used throughout the app — stripping Greek diacritics and Farsi harakat and uppercasing Latin script — so that a learner is not penalised for a guess that differs from the stored form only in diacritics their keyboard does not easily produce. Word length (5), attempt count (6), the permitted letter pattern, and the on-screen keyboard layout are all resolved per-language from `getWordleConfig` (Section C.6.1) rather than hard-coded, and the board and keyboard render right-to-left when the active language’s config marks it as such. The due-word and tracked-word sets are read through refs rather than directly inside the round-initialisation callback, specifically so that a mid-round `reviewWord` call — which updates `userVocabulary`, and so would otherwise change those sets’ identity — does not itself retrigger initialisation and silently restart the round before the learner reaches the collect step. Unlike Word of Wonders (Section D.5.2), Wordle does not use the self-graded Forged/Hard/Good/Easy prompt at all: guessing the secret word within the attempt limit always records a review with a fixed `GOOD` rating, whether or not that word was genuinely due, and running out of attempts records no review at all rather than a failing one. Coins are awarded through the same `userProgress` update path regardless of outcome (5 for a win, 1 for a loss).

D.5.2 Word of Wonders: Procedural Crossword Generation

D.5.2.1 Word Placement and Intersection Search

A level is built on an 8×8 grid, starting from the practice word and greedily placing further words from the learner’s tracked vocabulary wherever they intersect an already-placed word at a shared letter. A candidate placement is rejected if it would overlap a perpendicular word at a non-shared cell, or touch an existing word edge-to-edge without actually crossing it, which is what keeps every placed word legible as a separate word rather than fusing into an unintended run of letters. Algorithm 1 summarises the core loop, run until either every candidate word has been tried or `MAX_WORDS` (12) placements have been made.

Algorithm 1 Word of Wonders board generation (per level)

```

1: words  $\leftarrow$  learner's tracked vocabulary, filtered to valid length and script for the target language
2: practice  $\leftarrow$  due word if one exists and survives the filter above, else a random word from words
3: place practice horizontally at a random row on the empty board
4: sort remaining words by the priority heuristic of Algorithm 2 (descending)
5: for each word in sorted words, while placed count < MAX_WORDS do
6:   if placing word would exceed the letter budget (Section D.5.2, next) then
7:     skip word
8:   end if
9:   for each already-placed word p do
10:    for each shared-letter intersection between p and word do
11:      if the resulting placement is valid (in bounds, no illegal overlap) then
12:        place word at that intersection; break out of both loops
13:      end if
14:    end for
15:  end for
16: end for
17: trim the board to the bounding box of all placed cells

```

D.5.2.2 Letter-Budget and Word-Priority Heuristics

The player picks letters for each word from a single shared pool of physical letter tiles shown alongside the board, rather than a full keyboard, so that pool must stay small regardless of how many words end up on the board. `MAX_LETTERS_LENGTH` (8) bounds it: a candidate word is only placed if the *multiset union* of its letters with the letters already in use — not the sum of both counts — stays within that budget, so a word reusing letters already on the board is cheaper to add than one introducing entirely new letters. Candidate words are tried in priority order for exactly this reason (Algorithm 2): ranked first by how many letters they share with the pool already in use, and second (as a tiebreak) by how few new letters they would add, so the generator exhausts low-cost, letter-reusing words before it is forced to spend budget on words that would grow the pool.

Algorithm 2 Word priority for placement ordering

```

1: shared, added  $\leftarrow$  0, 0
2: for each letter  $\ell$  appearing in word or already in the pool do
3:   before  $\leftarrow$  pool's current count of  $\ell$ 
4:   after  $\leftarrow$   $\max(\textit{before}, \textit{word's count of } \ell)$ 
5:   if after > before then
6:     added  $\leftarrow$  added + (after − before) {word needs more of  $\ell$  than the pool has}
7:   else
8:     shared  $\leftarrow$  shared + before {pool already covers  $\ell$ ; credit its full existing count}
9:   end if
10: end for
11: return (shared, −added) {sort descending: more shared letters first, then fewer added letters}

```

D.5.2.3 Due-Word Injection as the Practice Word

The practice word is the one word in a Word of Wonders level whose discovery is treated as an FSRS review, and it is only ever marked as such — surfacing the self-graded rating prompt of Section C.5.2 on the finish screen — when it was actually drawn from the due-word queue rather than substituted at random because nothing was due. Every other word placed on the board is practice in the ordinary sense (the learner has tracked it before and may re-encounter it here) without being graded as a spaced-repetition event, keeping FSRS state changes limited to words genuinely

scheduled for review rather than every word that happens to appear on a given board. This due-word gate is specific to Word of Wonders; the Wordle-style game (Section D.5.1) grades every completed round regardless of due-word status, rather than gating on it the same way.

Chapter E

Evaluation & Results

Chapter F

Conclusions & Future work

F.1 Contributions

F.2 Limitations

The work presented in this thesis has several limitations that must be acknowledged honestly. **No empirical user study.** The pedagogical effectiveness of the system has not been evaluated in a controlled experiment. **Limited content library.** The current content library is limited in scope and diversity.

Steep onboarding barriers. The Glosy same of LingQ, incorporates real-world texts, its interface poses steep onboarding barriers for novice learners [23].

Lack of social scaffolding. The system lacks mechanisms for social scaffolding through peer or native-speaker interaction.

F.3 Future Work

The following directions are identified as high-value:

- **Automated content annotation pipeline:** A pipeline that ingests a raw video URL, transcribes speech using a large language model, aligns the transcript to video timestamps, tokenises and POS-tags the output, maps tokens to the word database, and produces a fully

annotated reel record—reducing the per-reel curation cost from hours to minutes.

- **Empirical evaluation:** A randomised controlled study measuring vocabulary gain, retention at one week and one month, and learner engagement metrics, comparing Glosy against a no-game control condition and against Duolingo.
- **Multi-language expansion:** The schema and service layer are language-agnostic (language is a foreign key throughout); adding a new language requires only content ingestion, not code changes.
- **Collaborative learning features:** Shared vocabulary challenges, leaderboards scoped to friend groups, and comment-driven vocabulary discussion on individual reels, supporting the social dimension of Vygotsky’s theory.
- **Vocabulary synchronisation performance:** Replacing the serial per-word UPDATE loop in `syncController.js` with a bulk `INSERT ... ON CONFLICT DO UPDATE` statement to support large offline sessions without latency degradation.

Bibliography

- [1] Hermann Ebbinghaus. *Memory: A Contribution to Experimental Psychology*. Teachers College, Columbia University, New York, 1885. Translated by Henry A. Ruger and Clara E. Bussenius, 1913.
- [2] Jaap M. J. Murre and Joeri Dros. Replication and analysis of ebbinghaus' forgetting curve. *PLOS ONE*, 10(7):e0120644, 2015.
- [3] Zirguez. Leitner system (alternative). Wikimedia Commons, https://commons.wikimedia.org/wiki/File:Leitner_system_alternative.svg, 2012. CC0 1.0 Public Domain Dedication. Accessed: 2026-09-15.
- [4] Kent State University. Global language learning trends and statistics. Kent State University Department of Modern & Classical Language Studies Blog, August 2025. Accessed: 2026-08-26.
- [5] Mordor Intelligence. Online language learning market size & share analysis - growth trends and forecast (2026 - 2031). Technical report, Mordor Intelligence, August 2026.
- [6] TalentLMS. The talentlms 2026 annual l&d benchmark report. Technical report, TalentLMS, 2026.
- [7] Neal Mohan. The future of YouTube in 2026: Annual letter from the CEO, January 2026.
- [8] Mark Zuckerberg. Meta's third quarter 2025 earnings conference call, October 2025.
- [9] Kepios. Essential TikTok statistics, audience reach and demographics, March 2025.
- [10] Similarweb. Digital 2026: Global overview report (android app insights), October 2025.
- [11] Pew Research Center. Teens, social media and AI chatbots 2025. Technical report, Pew Research Center, December 2025.

- [12] Pew Research Center. Americans' social media use 2025. Technical report, Pew Research Center, November 2025.
- [13] Hasan Cagli. Short-form video statistics 2026 (TikTok, reels & shorts), April 2026.
- [14] Wistia, Inc. State of video report: Video marketing statistics for 2026. Technical report, Wistia, April 2026.
- [15] Vidyard. Video in business benchmark report. Technical report, Vidyard, 2024.
- [16] Oxana Zelevskaya. Hyper-casual games: definition, genre history and specifics, 2022.
- [17] Anthony Pizzo. Hypercasual and hybrid-casual video gaming: A digital leisure perspective. *Leisure Sciences*, 47:1333–1352, 05 2023.
- [18] Business Research Insights. Hyper casual gaming market size, share, growth, trends, global industry analysis, by type (android games and ios games) by application (men users and women users), regional insights and forecast from 2026 to 2035. <https://www.businessresearchinsights.com/market-reports/hyper-casual-gaming-market-103060>, 2026.
- [19] Benjamin Paaßen, Thekla Morgenroth, and Michelle Stratemeyer. What is a true gamer? the male gamer stereotype and the marginalization of women in video game culture. *Sex Roles*, 76:421–435, 2017.
- [20] Kevin Anderton. Hypercasual female gamers are taking over the industry [infographic], September 2020. Accessed: 2026-08-27.
- [21] Alyssa Celatti. Study: Millions of us seniors aged 65+ spent \$500 on gaming in the past year, 2022. Accessed: 2026-08-27.
- [22] Business Research Insights. Language learning application market size, share, growth, and industry analysis. <https://www.businessresearchinsights.com/market-reports/language-learning-application-market-102456>, 2026. Accessed: 2026-07-28.
- [23] Athanasios Karasimos. The battle of language learning apps: a cross-platform overview. *Research Papers in Language Teaching and Learning*, 12(1):150–166, 2022.
- [24] Duolingo, Inc. shareholder letter q2 2026, 2026.

- [25] Duolingo, Inc. Duolingo launches 148 new language courses. Press Release, Duolingo Investor Relations, <https://investors.duolingo.com/news-releases/news-release-details/duolingo-launches-148-new-language-courses>, April 2025. Accessed: 2026-07-28.
- [26] Alex Gilmore. Authentic materials and authenticity in foreign language learning. *Language Teaching*, 40(2):97–118, 2007.
- [27] Hui Liu and Hu Dong. Authentic language models and authenticity in the foreign language classrooms. *Journal of Educational Research and Policies*, 6:61–66, 08 2024.
- [28] Sebastian Deterding, Dan Dixon, Rilla Khaled, and Lennart Nacke. From game design elements to gamefulness: Defining gamification. In *Proceedings of the 15th International Academic MindTrek Conference: Envisioning Future Media Environments*, pages 9–15. ACM, 2011.
- [29] Jan L. Plass, Bruce D. Homer, and Charles K. Kinzer. Foundations of game-based learning. *Educational Psychologist*, 50(4):258–283, 2015.
- [30] Lev S. Vygotsky. *Mind in Society: The Development of Higher Psychological Processes*. Harvard University Press, Cambridge, MA, 1978. Edited by Michael Cole, Vera John-Steiner, Sylvia Scribner, and Ellen Souberman.
- [31] James P. Lantolf, editor. *Sociocultural Theory and Second Language Learning*. Oxford University Press, Oxford, 2000.
- [32] HelloTalk Inc. HelloTalk – learn languages with native speakers. <https://www.hellotalk.com>, 2026. Accessed: 2026-09-15.
- [33] Stephen D. Krashen. *The Input Hypothesis: Issues and Implications*. Longman, London, 1985.
- [34] Marcella Hu and I.S.P. Nation. Unknown vocabulary density and reading comprehension. *Reading in a Foreign Language*, 13(1):403–430, 2000.
- [35] I.S.P. Nation. *Learning Vocabulary in Another Language*. Cambridge University Press, Cambridge, UK, 2001.
- [36] George A. Miller. The magical number seven, plus or minus two: Some limits on our capacity for processing information. *Psychological Review*, 63(2):81–97, 1956.

- [37] John Sweller. Cognitive load during problem solving: Effects on learning. *Cognitive Science*, 12(2):257–285, 1988.
- [38] Sebastian Leitner. *So lernt man lernen: Der Weg zum Erfolg*. Herder, Freiburg, 1972.
- [39] Burr Settles and Brendan Meeder. A trainable spaced repetition model for language learning. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 1848–1858, 2016.
- [40] Junyao Ye, Jingyong Su, and Yilong Cao. A stochastic shortest path algorithm for optimizing spaced repetition scheduling, 08 2022.
- [41] Open Spaced Repetition. The algorithm. GitHub wiki, <https://github.com/open-spaced-repetition/awesome-fsrs/wiki/The-Algorithm>, 2026. Accessed: 2026-09-16.
- [42] Henry L. Roediger and Jeffrey D. Karpicke. Test-enhanced learning: Taking memory tests improves long-term retention. *Psychological Science*, 17(3):249–255, 2006.
- [43] Robert Vanderplank. The value of teletext sub-titles in language learning. *ELT Journal*, 42(4):272–281, 1988.
- [44] Elke Peters and Stuart Webb. Incidental vocabulary acquisition through viewing L2 television and factors that affect learning. *Studies in Second Language Acquisition*, 40(3):551–577, 2018.
- [45] Maribel Montero Perez, Elke Peters, and Piet Desmet. Vocabulary learning through viewing video: The effect of two enhancement techniques. *Computer Assisted Language Learning*, 31(1-2):1–26, 2018.
- [46] Robert Godwin-Jones. Games in language learning: Opportunities and challenges. *Language Learning & Technology*, 18(2):9–19, 2014.
- [47] Julie M. Sykes and Jonathon Reinhardt. *Language at Play: Digital Games in Second and Foreign Language Teaching and Learning*. Pearson Education, New York, 2013.
- [48] Zichu Yang and Bowen Sun. Hyper-casual endless game based dynamic difficulty adjustment system for players replay ability. In *2020 IEEE Intl Conf on Parallel & Distributed Processing with Applications, Big Data & Cloud Computing, Sustainable Computing & Commu-*

- nications, *Social Computing & Networking (ISPA/BDCLOUD/SocialCom/SustainCom)*, pages 860–866. IEEE, 2020.
- [49] Gediminas Adomavicius and Alexander Tuzhilin. Toward the next generation of recommender systems: A survey of the state-of-the-art and possible extensions. *IEEE Transactions on Knowledge and Data Engineering*, 17(6):734–749, 2005.
- [50] Robin Burke. Hybrid recommender systems: Survey and experiments. *User Modeling and User-Adapted Interaction*, 12(4):331–370, 2002.
- [51] Francesco Ricci, Lior Rokach, and Bracha Shapira. Introduction to recommender systems handbook. In Francesco Ricci, Lior Rokach, Bracha Shapira, and Paul B. Kantor, editors, *Recommender Systems Handbook*, pages 1–35. Springer, 2011.
- [52] Aleksandra Klačnja-Milićević, Boban Vesin, Mirjana Ivanović, and Zoran Budimac. E-learning personalization based on hybrid recommendation strategy and learning style identification. *Computers & Education*, 56(3):885–899, 2011.

Abbreviations

AI	Artificial Intelligence
CLT	Cognitive Load Theory
FSRS	Free Spaced Repetition Scheduler
HLR	Half-Life Regression
L2	Second Language
MALL	Mobile-Assisted Language Learning
SLA	Second Language Acquisition
SRS	Spaced Repetition System
ZPD	Zone of Proximal Development

Glossary

Algorithm	A sequence of instructions for solving a problem
Comprehensible Input	Language input a learner can understand despite containing some unknown elements, per Krashen's Input Hypothesis ($i + 1$)
Forgetting Curve	The exponential decline of memory retention over time in the absence of reinforcement, first described by Ebbinghaus
Gamification	The addition of game-design elements (points, streaks, badges) to a non-game task, as distinct from game-based learning
Lexical Coverage	The proportion of word tokens in a text that a reader already knows; 98% coverage is the established threshold for unassisted comprehension
Recommender System	An algorithm that ranks or filters candidate items for a user based on content features, collaborative behaviour, or both
Spaced Repetition	A learning technique that schedules review of an item at increasing intervals timed to precede predicted memory decay
Testing Effect	The finding that actively retrieving information from memory produces better long-term retention than passive re-study

Zone of Proximal
Development

The gap between what a learner can do unaided and what they can do with appropriate scaffolding, per Vygotsky's sociocultural theory